(a) A real-feedback proxy

Each submission is scored on a new random draw: scores mix method quality with draw difficulty

↓

At budget 0 a pair's predictions are constant, so $r(1-r) = \text{Brier} - \text{ECE}^2$

↓

Eight reports yield 69 real pairs with recovered success rates r

↓

Score any label-only predictor on the same pairs

(b) What limits label-only prediction

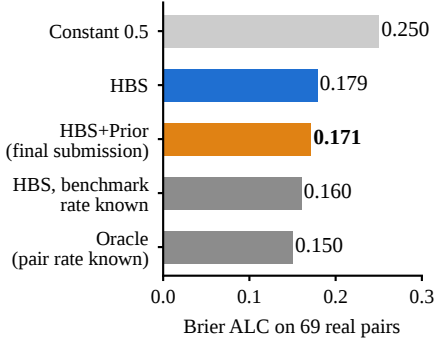

Unknown benchmark difficulty accounts for 65% of the gap to the oracle; the LLM benchmark prior closes 40% of it.