(a) A real-feedback proxy

Each submission is scored on a new random draw: scores mix method quality with draw difficulty

At budget 0 a pair's predictions are constant, so $r(1-r) = \text{Brier} - \text{ECE}^2$

Eight reports yield 69 real pairs with recovered success rates r

Score any label-only predictor on the same pairs
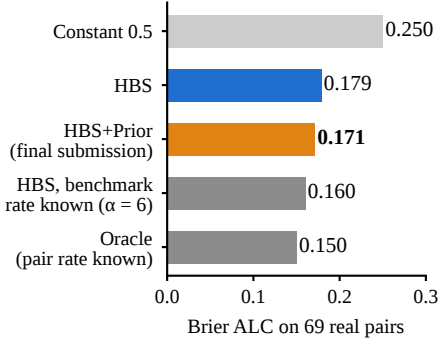
(b) What limits label-only prediction